

Bootstrap-t Comparison

Mario Ibanez

March 26, 2016

Question

Let x_1, x_2, \dots, x_{20} each be independent values from an exponential distribution with mean $= 1/\lambda$. (In other words, $f(x) = \lambda e^{-\lambda x}$). Use the Bootstrap-t method to construct a 95% confidence interval for the parameter λ .

Answer

Below is a random sample of 20 values from the exponential distribution:

```
# Set seed
set.seed(1234)

# Load packages and data
library(ggplot2)
library(knitr)

# 20 random exponential values with mean = 1/10
lambda <- 10
data <- rexp(n = 20, rate = lambda); data
```

```
## [1] 0.2501758605 0.0246758883 0.0006581957 0.1742746090 0.0387182584
## [6] 0.0089949671 0.0824081515 0.0202617901 0.0838040319 0.0760430301
## [11] 0.1880076678 0.1596105418 0.1658662384 0.3052458100 0.1750680133
## [16] 0.0031725529 0.0876960574 0.0014613740 0.1835064027 0.0519341271
```

The 95% confidence interval will be calculated two ways.

Method 1

This method will use the delta method to estimate the variance of $\hat{\theta} = 1/\bar{x}$. We have:

$$S_{\hat{\theta}} = \frac{S}{\bar{x}^2 \sqrt{n}}$$
$$\hat{\theta} = \frac{1}{\bar{x}}$$

Below are the values of $\hat{\theta}$ and $S_{\hat{\theta}}$

```
# Values for original sample
theta_hat <- 1 / mean(data); theta_hat
```

```
## [1] 9.60807
```

```
s_theta_hat <- sd(data) / (mean(data)^2 * sqrt(length(data))); s_theta_hat
```

```
## [1] 1.862861
```

Now we can bootstrap 1000 samples from this data. First, this is the function that will be used to generate the “t-statistic” for each bootstrapped sample:

```
# Function to calculate statistics
statistic_generator <- function(x, theta_hat){
  theta_hat_star <- 1 / mean(x)
  s_theta_hat_star <- sd(x) / (mean(x)^2 * sqrt(length(x)))
  return((theta_hat_star - theta_hat)/s_theta_hat_star)
}
```

This function will calculate the values $\hat{\theta}$, $S_{\hat{\theta}}$, and T_i^* when given a bootstrapped sample.

Now 1000 bootstrapped samples are calculated:

```
boot_strapped_data <- matrix(data = sample(x = data, size = 20*1000, replace = TRUE),
                             ncol = 20, byrow = TRUE)
```

Now the statistics are calculated:

```
boot_statistics <- matrix(apply(X = boot_strapped_data, MARGIN = 1,
                                FUN = statistic_generator, theta_hat = theta_hat),
                           ncol = 1, byrow = TRUE)
```

Now we can calculate the 95% confidence interval using the 2.5th and 97.5th percentiles of the “t-statistics”:

```
sorted_t_statistics <- sort(boot_statistics[, 1])
lower_limit <- theta_hat - sorted_t_statistics[1000*0.975] * s_theta_hat; lower_limit
```

```
## [1] 6.98008
```

```
upper_limit <- theta_hat - sorted_t_statistics[1000*0.025] * s_theta_hat; upper_limit
```

```
## [1] 15.10695
```

Now we can try to run this same method 100 times in order to approximate the true coverage probability of this method.

```
results_vector <- rep(0, 100)
confidence_width <- rep(0, 100)

for(i in 1:100){
  boot_strapped_data <- matrix(data = sample(x = data, size = 20*1000, replace = TRUE),
                               ncol = 20, byrow = TRUE)
```

```

boot_statistics <- matrix(apply(X = boot_strapped_data, MARGIN = 1,
                             FUN = statistic_generator, theta_hat = theta_hat),
                        ncol = 1, byrow = TRUE)

sorted_t_statistics <- sort(boot_statistics[, 1])
lower_limit <- theta_hat - sorted_t_statistics[1000*0.975] * s_theta_hat; lower_limit
upper_limit <- theta_hat - sorted_t_statistics[1000*0.025] * s_theta_hat; upper_limit

results_vector[i] <- ((lambda < upper_limit) && (lambda > lower_limit))
confidence_width[i] <- upper_limit - lower_limit
}

```

Out of 100 simulations, 100 contained the true value of lambda. The average confidence interval width was 7.7427197

Method 2

This method creates a confidence interval instead for the value $1/\lambda$ directly, in other words, $\hat{\theta} = \bar{x}$. At the end, the confidence interval is transformed into a confidence interval for λ

```

# Values for original sample
theta_hat2 <- mean(data); theta_hat2

```

```
## [1] 0.1040792
```

```
s_theta_hat2 <- sd(data) / sqrt(length(data)); s_theta_hat2
```

```
## [1] 0.0201794
```

Now we can bootstrap 1000 samples from this data. First, this is the function that will be used to generate the “t-statistic” for each bootstrapped sample:

```

# Function to calculate statistics
statistic_generator2 <- function(x, theta_hat2){
  theta_hat_star <- mean(x)
  s_theta_hat_star <- sd(x) / sqrt(length(x))
  return((theta_hat_star - theta_hat2)/s_theta_hat_star)
}

```

So that the comparison is more equivalent, the same bootstrapped data will be used. Of course new “t-statistics” need to be calculated however.

```

boot_statistics2 <- matrix(apply(X = boot_strapped_data, MARGIN = 1,
                              FUN = statistic_generator2, theta_hat = theta_hat2),
                          ncol = 1, byrow = TRUE)

```

Now we can calculate the 95% confidence interval using the 2.5th and 97.5th percentiles of the “t-statistics”:

```
sorted_t_statistics2 <- sort(boot_statistics2[, 1])
lower_limit2 <- theta_hat2 - sorted_t_statistics2[1000*0.975] * s_theta_hat2; lower_limit2
```

```
## [1] 0.0642807
```

```
upper_limit2 <- theta_hat2 - sorted_t_statistics2[1000*0.025] * s_theta_hat2; upper_limit2
```

```
## [1] 0.1521753
```

These lower and upper limits are for the value of $1/\lambda$.

Now we can try to run this same method 100 times in order to approximate the true coverage probability of this method.

```
results_vector2 <- rep(0, 100)
confidence_width2 <- rep(0, 100)

for(i in 1:100){
  boot_strapped_data <- matrix(data = sample(x = data, size = 20*1000, replace = TRUE),
                               ncol = 20, byrow = TRUE)

  boot_statistics <- matrix(apply(X = boot_strapped_data, MARGIN = 1,
                                FUN = statistic_generator2, theta_hat = theta_hat2),
                            ncol = 1, byrow = TRUE)

  sorted_t_statistics <- sort(boot_statistics[, 1])
  upper_limit <- 1 / (theta_hat2 - sorted_t_statistics[1000*0.975] * s_theta_hat2)
  lower_limit <- 1 / (theta_hat2 - sorted_t_statistics[1000*0.025] * s_theta_hat2)

  results_vector2[i] <- ((lambda < upper_limit) && (lambda > lower_limit))
  confidence_width2[i] <- upper_limit - lower_limit
}
```

Out of 100 simulations, 100 contained the true value of λ . The average confidence interval with was 8.6291913

This method also had good performance but the confidence intervals were much wider on average.