

Statistical Computing HW 4

Mario Ibanez

January 30, 2016

Problem 4.3)

Give an efficient algorithm to simulate the value of a random variable X such that

$$P\{X = 1\} = 0.3$$

$$P\{X = 2\} = 0.2$$

$$P\{X = 3\} = 0.35$$

$$P\{X = 4\} = 0.15$$

Answer

Derivation

The CDF of the distribution is:

$$F(x) = \begin{cases} 0 & x < 1 \\ 0.3 & 1 \leq x < 2 \\ 0.5 & 2 \leq x < 3 \\ 0.85 & 3 \leq x < 4 \\ 1 & 4 \leq x \end{cases}$$

This can be used along with a standard uniform random variable to generate values of the random variable X .

Algorithm

This is slightly different than the inverse transform method in the book but the outcome is the same.

1. Generate a number u from the standard uniform distribution
2. Set $x = F^{-1}(u) + 1$
3. Repeat n times

Program

Below is a program and histogram using the derivation and algorithm. 10,000 values were generated. The distribution is quite good.


```
# F inverse function
F_inv <- function(u){
  x <- rep(0, length(u))
  if(u < 0.30)
    x <- 0 + 1
  else if((u >= 0.30) && (u < 0.50))
    x <- 1 + 1
  else if((u >= 0.50) && (u < 0.85))
    x <- 2 + 1
  else if((u >= 0.85) && (u < 1))
```

```

    x <- 3 + 1
    return(x)
}
# This line is generates the x values
x <- sapply(array(runif(10^4)), FUN = F_inv)

```

Distribution of 10,000 generated X values



Problem 4.14a)

Let X be a binomial random variable with parameters n and p . Suppose that we want to generate a random variable Y whose probability mass function is the same as the conditional mass function of X given that $X \geq k$, for some $k \leq n$. Let $\alpha = P\{X \geq k\}$ and suppose that the value of α has been computed. Give the inverse transform method for generating Y .

Answer

Derivation

Algorithm

Program

Analytical Result

Problem 4.15)

Give a method for simulating X , having the probability mass function p_j , $j = 5, 6, 7, \dots, 14$, where

$$p_j = \begin{cases} 0.11 & \text{when } j \text{ is odd and } 5 \leq j \leq 13 \\ 0.09 & \text{when } j \text{ is even and } 6 \leq j \leq 14 \end{cases}$$

Use the text's random number sequence to generate X .

Answer

(What is the text's random number sequence?)

The method used will be the composition approach.

Derivation

The composition approach works by finding α , $p_j^{(1)}$, and $p_j^{(2)}$ so that

$$P\{X = j\} = \alpha p_j^{(1)} + (1 - \alpha) p_j^{(2)}$$

Let $\alpha = 0.55$ and let

$$p_j^{(1)} = P\{X = j\} = \begin{cases} 0 & \text{for } x \text{ even} \\ 0.2 & \text{for } x \text{ odd} \end{cases}$$

and

$$p_j^{(2)} = P\{X = j\} = \begin{cases} 0 & \text{for } x \text{ odd} \\ 0.2 & \text{for } x \text{ even} \end{cases}$$

In other words, if j_0 is odd, then $P(X = j_0) = (0.55)(0.20) + (0.45)(0) = 0.11$ and if j_0 is even then $P(X = j_0) = (0.55)(0) + (0.45)(0.20) = 0.09$. This is the desired outcome.

Algorithm


1. Generate u_1 from standard uniform distribution
2. Generate u_2 from standard uniform distribution
3. If $u_1 < \alpha$, (odd case) set $x = 2(\text{floor}(5u_2)) + 5$
4. Otherwise (even case) set $x = 2(\text{floor}(5u_2)) + 6$
5. Go back to step 1, repeat n times

Program

```

# Initialize variables, generate uniform random numbers
N <- 10^4
alpha <- 0.55
uniform1 <- runif(N)
uniform2 <- runif(N)
x_dist <- c(rep(0, N))
# These two lines generate the x values
x_dist[uniform1 < 0.55] <- 2*floor(5*uniform2[uniform1 < 0.55])+5
x_dist[x_dist == 0] <- 2*floor(5*uniform2[uniform1 >= 0.55])+6

```



Problem 5.10) A casualty insurance company has 1000 policyholders, each of whom will independently present a claim in the next month with probability 0.05. Assuming that the amounts of the claims made are independent exponential random variables with mean \$800, use simulation to estimate the probability that the sum of these claims exceeds \$50,000.

Answer

Derivation

In order to generate random values from an exponential distribution with mean 800, we'll use the fact that the CDF of this distribution is

$$F(x) = 1 - e^{-x/800}$$

Given that $F(x)$ has a standard uniform distribution, then:

$$\begin{aligned}
u &= 1 - e^{-x/800} \\
e^{-x/800} &= 1 - u \\
\frac{-x}{800} &= \ln(1 - u) \\
x &= (-800)\ln(u)
\end{aligned}$$

Note that the random variables U and $1 - U$ have the same distribution if U is standard uniform.

Also, since each policyholder has a 5% chance of making a claim, then the number of claims made out of 1000 has a binomial distribution with $n = 1000$ and $p = 0.05$. The random variable we are interested in is then

$$S = X_1 + X_2 + X_3 + \dots + X_N$$

where the X_i are *iid* exponential with mean 800 and N is binomial with $n = 1000$ and $p = 0.05$. For simplicity in the program, values from the binomial distribution will be done using an R function, but the exponential values will be generated using values from a uniform distribution according to the derivation above.

Algorithm

- For each trial:
 - Generate a value n from Binomial(1000, 0.05)
 - Generate n values u_1, u_2, \dots, u_n from $U(0,1)$
 - Calculate each of the n claim amounts x by evaluating $x_i = (-800)\ln(u_i)$
 - Find the sum of the n claims for that trial
- After the sum of each trial is calculated, then determine how many are greater than \$50,000

The R code appears different than this algorithm, but in practice this is what is happening. The R language benefits from avoiding explicitly writing *for* loops, and instead using other methods like *apply*, *lapply*, and *sapply*.

Program

The program uses the *generate_claims()* function to generate n claims from an exponential distribution, where n has a binomial distribution. 100,000 trials are run, and then the sum of the claims from each trial is calculated.

```

# Set seed, set N = number of simulations
set.seed(1234)
N = 10^5
# Returns variable number claims from exponential distribution
generate_claims <- function(vector, holders = 1000, chance = 0.05, beta = 800){
  # Number of claims that will occur is binomial(holders, chance)
  num_of_claims <- rbinom(1, holders, chance)
  # Generate claim amounts from exponential distribution for each claimant
  claims <- -beta * log(runif(num_of_claims))
  # Return vector of claims
  return(claims)
}

```


```

# Initialize claims list
claims <- rep(list(rep(0,100)), N)
# Simulate N times
claims <- lapply(claims, generate_claims)
# Find the sum for each trial
distribution <- sapply(claims, sum)

```

Below is a histogram of the results. The distribution is a sum of exponential random variables and thus is a continuous distribution. It appears to be slightly right tailed with a mean around \$40,000.

Distribution of Claim Amounts



The following R command returns the estimated probability that the sum of the claims is above \$50,000:

```
mean(distribution>50000)
```

```
## [1] 0.10661
```

Out of 100,000 trials, approximately 0.10661 trials resulted in an outcome where the total value of the claims exceeded \$50,000.

Analytical Result

find out how to do this analytically